

Decksmith

AI that generates presentation decks you can **actually edit** — native .pptx and .key files that are designed, not flattened into images.

The deck AI hands you isn't the deck you can send.

A salesperson prepping a client pitch — under time pressure — wants a deck that's **ready**, not a draft to rebuild.

PAIN · P1

Looks templated

Generic, off-brand output that reads "made by AI" at a glance.

PAIN · P2

Wrong structure

Passable visuals wrapped around the wrong argument and framing.

PAIN · THE KILLER

Can't edit it

Export lands as flattened images — so you redo the whole thing by hand.

The time it "saves" comes straight back as rework.

Every AI deck tool is web-first.
The file you download breaks.


30–40%

of slides get **flattened to images**
when Gamma exports to .pptx.
Incumbents polish the browser view;
the moment you need PowerPoint or
Keynote, fidelity collapses.

sourced: [market-research](#)

The gap = native fidelity – unclaimed ground.

Generate a deck. Export a file where every object is real.



EDITABLE

Native text, shapes and charts — not a picture of a slide.

"A designed client deck you can actually edit in PowerPoint."

DESIGNED

On-brand and composed, not a visible template.

YOURS

Open it, change it, keep working — no lock-in to a viewer.

BOTTOM-UP • WIDE ERROR BARS

~\$750M SAM assumption

Obtainable market for editable, designed decks
— sized from reachable segments, not a top-down share of "all presentations."

WHO'S ALREADY HERE

Gamma

Microsoft Copilot

Canva

Beautiful.ai

Pitch

Tome

The threat: Copilot and Canva are entering the same wedge. The moat has to be the fidelity + design **engine**, not the app around it.

Win one arena first: salespeople under deadline.

Beachhead

Salespeople prepping a client pitch under time pressure — a sharp, reachable job.

Positioning

"Actually-editable + designed native decks" — a wedge strong enough to beat the PowerPoint habit.

Pricing

\$15–25 / mo standalone. Open willingness-to-pay — a hypothesis to test, not a set price.

Channels

Sales-enablement communities + product-led sharing of the exported deck.

One metric decides whether the concept is real.

Edit-fidelity	$\geq 90\%$ of exported objects natively editable (M-edit-fidelity) — the concept-proving gate.
North Star	Kept decks per week — the value that has to repeat.
Guardrail	$\text{COGS} \leq \$1 / \text{deck}$ — LLM inference is a cost line, not overhead.
Red line	No flatten-to-image shortcut — it would fake the entire wedge.

A 6-week Concept Test with one learning goal.

Can we build native edit-fidelity at scale? (H-001) Everything else is sized to not over-invest before that signal.

DEV · F-1 Native export slice Generate → export one real .pptx; measure fidelity.	GTM · A-1 Demand probe Seed the wedge in one sales community; count qualified trials.	BACK-OFFICE · T-1 Fidelity harness Repeatable, honest edit-fidelity measurement + event capture.
≥ 90% Validated — build on it.	70–90% Inconclusive — iterate the slice.	< 70% Refuted — rethink the engine or pivot.

THE DECISION ON THE TABLE

Fund a 6-week Concept Test to buy one honest signal.

Commit	~6 weeks · founder + 1 eng · no paid acquisition — the smallest test that can answer it.
Decide	Go / no-go on the native edit-fidelity read: $\geq 90\%$ → build on · $< 70\%$ → stop or pivot.
Today	Everything is [assumption] — this is the cheapest way to get the first real evidence.

If the export is real, the wedge is real.